

A Comparative Study of Convolutional and Transformer Architectures Under Adversarial Perturbations: Gradient Characteristics and Transferability Analysis

Çağrı Şahin and Özal Yıldırım

Abstract—How convolutional neural networks (CNNs) and Vision Transformers (ViTs) differ under adversarial perturbations remains an open question. This paper presents a behavioral comparison of the two architecture families under a matched adversarial objective and threat model, pairing robustness scores with gradient- and feature-level analyses. Using CIFAR-10 with adversarially trained ResNet-18 and ViT-Tiny models evaluated under AutoAttack at the standard budget of 8/255 on the full test set, we report four findings. First, the CNN attains higher absolute robust accuracy (35.74 versus 32.94 percent; McNemar p below 0.001), but a conditional decomposition shows this headline is composition-dependent: the PGD-10 gap traces entirely to the clean-accuracy gap (conditioned fooling 52.2 versus 52.3 percent), under AutoAttack the per-model conditional ordering slightly favors the ViT, yet on the jointly solved subset the CNN retains a significant edge under both attacks. Second, the apparent cross-architecture transfer asymmetry depends strongly on the conditioning protocol: raw rates manufacture an 8.3-point asymmetry from unequal clean accuracy, target-conditioned rates are statistically equivalent, while both-correct and successful-source protocols reveal a significant 5.3-point CNN-to-ViT advantage—so protocol choice can create, erase, or reverse directional conclusions. Third, scale-invariant sparsity measures show CNN gradients concentrate sensitivity more strongly (Hoyer 0.474 versus 0.449, paired tests p below 0.001), while cross-sample gradient alignment is higher for the ViT in absolute cosine (0.052 versus 0.038) but near zero in signed cosine for both. Fourth, ViT features degrade with an early-drop, mid-network-plateau profile, unlike the CNN’s monotonic depth profile. These results argue for protocol-explicit, multi-view robustness comparisons.

Index Terms—Adversarial robustness, convolutional neural networks, deep learning security, transfer attacks, Vision Transformer

I. INTRODUCTION

Deep neural networks have achieved remarkable success across various computer vision tasks, including image classification, object detection, and semantic segmentation. However, these models are vulnerable to adversarial examples—carefully crafted perturbations, often barely perceptible to humans, that can cause misclassification with high confidence [1], [2]. This vulnerability poses significant security concerns for safety-critical applications such as autonomous driving, medical diagnosis, and security systems.

Two dominant architectural paradigms have emerged in computer vision: Convolutional Neural Networks (CNNs) and Vision Transformers (ViTs). CNNs leverage local connectivity and weight sharing through convolutional operations, capturing hierarchical spatial features effectively [3]. In contrast, ViTs apply the self-attention mechanism [4] to image patches, enabling global context modeling from early layers [5]. While both architectures achieve competitive performance on clean data, their behavior under adversarial attacks exhibits fundamental differences that are not yet fully understood.

Previous research has extensively studied adversarial robustness for CNNs, leading to effective defense mechanisms such as adversarial training [6] and TRADES [7]. However, the comparative analysis of CNN and ViT robustness remains limited, with existing studies often focusing on *whether* one architecture is more robust than another, rather than *why* they exhibit different behaviors. Understanding the behavioral factors underlying these differences is crucial for developing more secure deep learning systems.

In this work, we move beyond simple robustness comparisons to characterize how CNNs and ViTs respond differently to adversarial perturbations. Our investigation centers on several interconnected questions: How do adversarial perturbations transfer between architectures, and how does the apparent directionality depend on the conditioning protocol used to measure it? Do gradient characteristics correlate with the observed robustness differences? How do adversarial perturbations propagate through Vision Transformer layers, and what feature-level dynamics lead to misclassification? What practical guidance for evaluating and deploying mixed-architecture ensembles follows from these findings?

This paper makes several contributions to understanding architectural differences in adversarial vulnerability. We provide a robustness evaluation using AutoAttack [8], a standardized strong benchmark, comparing adversarially trained ResNet-18 and ViT-Tiny under $\epsilon = 8/255$ L_∞ perturbations, and we decompose the absolute scores into clean-accuracy and conditional-susceptibility components, showing that the headline comparison is view-dependent. Our transfer attack analysis characterizes cross-architecture transferability under multiple established conditioning protocols, showing that the measured asymmetry is protocol-dependent: the raw metric manufactures an 8-point asymmetry from unequal clean ac-

The authors are with the Department of Software Engineering, Firat University, 23119 Elazığ, Turkey.

Corresponding author: Çağrı Şahin (e-mail: csahin@firat.edu.tr).

curacy, the target-conditioned metric shows statistical parity, and the stricter both-correct protocol reveals a significant CNN-to-ViT advantage on jointly solved samples—so protocol choice can create, erase, or reverse directional conclusions. We find that CNN gradients concentrate sensitivity more strongly in fewer input components under scale-invariant sparsity measures—a difference that survives a native-resolution ViT control and per-sample paired tests—while cross-sample gradient alignment is low for both architectures, moderately higher for the ViT in absolute cosine, and near zero in signed cosine for both. Furthermore, complementing earlier layer-wise robustness studies of ImageNet-scale ViTs [9], [10], we present a systematic layer-wise analysis of how adversarial perturbations affect the feature representations of an *adversarially trained* ViT, revealing an early-drop, mid-network-plateau degradation profile in which most directional deviation accumulates by mid-network rather than growing monotonically with depth. To support reproducible research, our complete experimental pipeline, trained models, and analysis tools will be made publicly available upon acceptance (see Data and Code Availability).

We begin by reviewing related work on adversarial attacks, defenses, and architecture-specific robustness studies in Section II. Section III describes our experimental methodology, and Section IV presents results across four complementary analyses: robustness comparison, transfer attacks, gradient characteristics, and feature degradation. We discuss the implications of our findings in Section V and conclude in Section VI.

II. RELATED WORK

A. Adversarial Attacks

Adversarial attacks aim to find minimal perturbations that cause misclassification. The Fast Gradient Sign Method (FGSM) [2] generates perturbations in a single step by following the sign of the loss gradient, with the perturbation magnitude bounded under the L_∞ norm. Projected Gradient Descent (PGD) [6] extends FGSM through iterative optimization, taking multiple gradient-sign steps and projecting back onto the ϵ -ball after each step (formal definitions in Section III-C).

The Carlini & Wagner (C&W) attack [11] formulates adversarial example generation as an optimization problem, often achieving higher success rates but at increased computational cost. More recently, AutoAttack [8] combines multiple attack strategies (APGD-CE, APGD-T, FAB-T, Square Attack) to provide a reliable robustness evaluation, becoming the de facto standard for benchmarking.

B. Defense Mechanisms

Adversarial training (AT) [6] remains the most effective defense, training models on adversarial examples generated during each batch:

$$\min_{\theta} \mathbb{E}_{(x,y) \sim \mathcal{D}} \left[\max_{\|\delta\|_\infty \leq \epsilon} \mathcal{L}(f_{\theta}(x + \delta), y) \right] \quad (1)$$

TRADES [7] provides a theoretically-grounded trade-off between natural accuracy and robustness:

$$\mathcal{L}_{\text{TRADES}} = \mathcal{L}_{CE}(f(x), y) + \beta \cdot \text{KL}(f(x) \| f(x_{adv})) \quad (2)$$

where β controls the trade-off parameter.

MART [12] addresses the issue of misclassified examples by emphasizing them during training. Other approaches include input transformation [13], certified defenses via randomized smoothing [14], and ensemble methods [15]. Recent advances in certified robustness include Adaptive Randomized Smoothing [16], which improves certification radii through multi-step defenses.

C. CNN Robustness

Extensive research has characterized the adversarial robustness of CNN architectures. Madry et al. [6] demonstrated that adversarial training with PGD attacks provides strong empirical robustness. The RobustBench leaderboard [17] tracks state-of-the-art CNN robustness, with WideResNet variants achieving the highest robust accuracy on CIFAR-10.

Studies have shown that model capacity correlates with achievable robustness [6], with wider and deeper networks demonstrating improved robust accuracy. Architecture-specific findings include the role of skip connections in adversarial transferability [18] and the distinct batch-normalization statistics that adversarial training induces at scale [19]; relatedly, auxiliary batch-normalization branches allow adversarial examples to improve clean accuracy [20]. Adversarial training is also known to reshape input-gradient structure itself, yielding more perceptually aligned [21] and sparser [22] gradients.

D. Vision Transformer Robustness

Vision Transformers [5], building on the self-attention mechanism introduced by Vaswani et al. [4], process images as sequences of patches; their internal representations differ measurably from CNNs' layer by layer [23]. Initial studies suggested that ViTs exhibit some inherent robustness properties due to their global receptive field [9]. Bai et al. [24] systematically investigated whether transformers are more robust than CNNs, finding that under fair training recipes transformers are not more adversarially robust than CNNs, while exhibiting superior robustness under distribution shift attributed to self-attention; Benz et al. [25] reported an early direct comparison including MLP-Mixer and linked the differences to frequency preferences. Our work extends this line of inquiry by analyzing adversarially trained models of comparable scale and providing behavioral characterizations through gradient characteristics and feature degradation analysis. Subsequent work revealed that ViTs remain vulnerable to adversarial attacks, particularly gradient-based methods [26], with patch-targeted attacks exploiting the attention mechanism directly [27] and attention shown to shift measurably under patch perturbations [28].

Shao et al. [10] demonstrated that adversarial training can improve ViT robustness—and that ViT features rely less on high-frequency content—though challenges remain due to training instability. ViT robustness to common corruptions and

occlusions has also been characterized [29]. More recently, Jain et al. [30] proposed Adaptive Attention Scaling Adversarial Training (AAS-AT) to address ViT-specific training challenges, achieving state-of-the-art robustness on CIFAR-10 and ImageNet-100; with diffusion-augmented adversarial training, ViTs have even been reported to surpass WideResNets on small-scale datasets [31]. Zhu et al. [32] strengthened transferable attacks on ViTs through gradient normalization scaling and high-frequency adaptation. Despite these advances, comprehensive behavioral comparisons between CNN and ViT robustness under unified evaluation frameworks remain limited.

E. Transfer Attacks

Adversarial examples often transfer between models, enabling black-box attacks [33]. The transferability depends on factors including model architecture, training data, and attack methods [34]. Foundational transfer studies already evaluate transferability on correctly classified samples [35], and standard attack papers condition their success rates for the same reason [36]; recent evaluation guidelines make such protocol choices explicit [37]. In ensemble design, Ravikumar et al. [38] restrict transfer measurement to samples correctly classified by both models precisely to account for accuracy differences between them—the same confound our conditioned metric controls.

Cross-architecture transfer between CNNs and ViTs has received increasing attention. Mahmood et al. [26] measured cross-architecture transfer with a both-correct protocol and found notably low CNN $\leftrightarrow$ ViT transferability; Naseer et al. [39] showed that adversarial examples crafted on ViTs transfer poorly under conventional attacks and proposed self-ensemble and token refinement to improve ViT-sourced transferability, and Wei et al. [40] improved ViT-sourced transfer by bypassing attention gradients and patch-wise dropout. Building on these, Token Gradient Regularization (TGR) [41] crafts transferable adversarial samples by exploiting ViT’s attention components, while Momentum Integrated Gradients (MIG) [42] achieves improved transferability across both ViTs and CNNs. Chen et al. [43] systematically analyze the factors affecting cross-architecture transferability, finding that architectural inductive biases—CNNs preferring high-frequency information while ViTs favor low-frequency patterns—significantly impact transfer success. Understanding these transfer patterns is crucial for developing robust ensemble systems and evaluating real-world attack scenarios.

F. Research Gap

While individual studies have examined CNN and ViT robustness separately, and the pitfalls of unfair CNN–transformer comparisons have been highlighted before [24], [44], existing comparative work typically reports robustness scores in isolation—often with inconsistent perturbation budgets, attack configurations, or model scales—rather than pairing the comparison with behavioral analysis. A recent study in this journal compares ViT and CNN adversarial robustness directly [45]; our work differs by evaluating on the full

test set with AutoAttack and paired statistics, by controlling transfer measurement for clean-accuracy differences, and by accompanying the comparison with gradient- and feature-level behavioral analyses.

Our work also engages the measurement question directly. Conditioning transfer evaluation on correctly classified samples is established practice [26], [35], [36], [38]; our contribution is to show, in a controlled same-experiment comparison, that the conditioning protocol itself is a first-order methodological variable: the unconditioned raw metric manufactures an 8.3-point spurious asymmetry when targets differ in clean accuracy, the target-correct metric shows parity, and the stricter both-correct and successful-source protocols reveal a genuine directional difference on the same data. We investigate gradient characteristics alongside transfer behavior, analyze layer-wise feature degradation in ViTs under adversarial perturbations, and connect these findings to practical guidance for evaluating and deploying mixed-architecture ensembles.

III. METHODOLOGY

A. Threat Model

We consider the standard L_∞ threat model where an adversary can perturb input images within a bounded region:

$$\mathcal{B}_\epsilon(x) = \{x' : \|x' - x\|_\infty \leq \epsilon\} \quad (3)$$

Following established conventions for CIFAR-10 [17], [46], we use $\epsilon = 8/255 \approx 0.031$ as the primary perturbation budget, with additional full-test PGD-10 evaluations at $\epsilon \in \{2/255, 4/255, 16/255\}$ for the epsilon sweep analysis.

We evaluate under both white-box (full model access) and black-box (transfer attack) scenarios. For white-box attacks, the adversary has complete knowledge of model architecture and parameters. For transfer attacks, adversarial examples are generated on a source model and evaluated on a different target model.

B. Model Architectures

a) *CNN Models: ResNet-18* [3]: A widely-used baseline CNN with residual connections, containing 11.2M parameters. We use the standard architecture designed for CIFAR-10 (initial 3×3 convolution instead of 7×7).

WideResNet-28-10 [47]: A wider variant of ResNet with 36.5M parameters, representing the state-of-the-art architecture for adversarial robustness on CIFAR-10. We use the pretrained robust checkpoint of Goyal et al. [48] distributed via RobustBench [17]; note that this model was trained with additional unlabeled data and a stronger recipe than our standard adversarial training, so it serves as a reference point rather than a matched comparison.

b) *Vision Transformer Models: ViT-Tiny*: We use the ViT-Tiny architecture from the `timm` library [49], adapted for CIFAR-10 by bilinearly upsampling inputs to 224×224 . The architecture uses 16×16 patches (196 patches per image), 192-dimensional embeddings, 12 transformer blocks with 3 attention heads each, MLP ratio of 4, and totals 5.7M parameters. The model is trained *from scratch* on CIFAR-10 (no ImageNet pretraining) with basic augmentation (random crop

and horizontal flip), then adversarially fine-tuned. Since ViT robustness on small datasets is known to be sensitive to the training recipe [50], [51], our architecture-level conclusions should be interpreted as holding *under this baseline training recipe* rather than for optimally trained ViTs.

C. Attack Methods

a) *FGSM*: Single-step attack with step size equal to the perturbation budget:

$$x_{adv} = x + \epsilon \cdot \text{sign}(\nabla_x \mathcal{L}_{CE}(f(x), y)) \quad (4)$$

b) *PGD*: Iterative attack with step size $\alpha = 2/255$ and random initialization:

$$x^{t+1} = \Pi_{\mathcal{B}_\epsilon(x)}(x^t + \alpha \cdot \text{sign}(\nabla_{x^t} \mathcal{L}_{CE}(f(x^t), y))) \quad (5)$$

with $x^0 = x + \mathcal{U}(-\epsilon, \epsilon)$. We use PGD-10 (10 iterations) for both adversarial training and evaluation throughout this work (Table I).

c) *AutoAttack*: We use AutoAttack [8], the standardized ensemble attack combining APGD-CE (Auto-PGD with cross-entropy loss), APGD-T (Auto-PGD with targeted DLR loss), FAB-T (targeted Fast Adaptive Boundary attack), and Square Attack (score-based black-box attack). This ensemble is a standardized strong benchmark for robustness evaluation.

D. Defense Methods

a) *Adversarial Training (AT)*: We employ a mixed adversarial training objective that combines clean and adversarial losses:

$$\mathcal{L} = (1 - \lambda) \mathcal{L}_{CE}(f_\theta(x), y) + \lambda \mathcal{L}_{CE}(f_\theta(x_{adv}), y) \quad (6)$$

where $\lambda = 0.5$ balances standard classification performance with adversarial robustness. The inner maximization generates adversarial examples using PGD-10 ($\alpha = 2/255$, $\epsilon = 8/255$), computed with the model in training mode so that attack generation and the parameter update see consistent batch-normalization statistics, following standard adversarial-training implementations [52]. This mixed objective, widely adopted in practice, helps maintain clean accuracy while building robustness, as training solely on adversarial examples can degrade natural performance. While alternative formulations such as TRADES [7] offer theoretically grounded trade-off mechanisms, this work focuses on standard adversarial training to isolate architectural effects from defense-specific interactions.

E. Training Protocol

a) *Clean Training*: For CNNs, we use SGD optimizer with initial learning rate 0.1, cosine annealing schedule [53], weight decay 5×10^{-4} , and 200 epochs. For ViTs, we use AdamW optimizer [54] with initial learning rate 10^{-3} , cosine annealing, weight decay 0.05, and 200 epochs.

b) *Adversarial Training*: Starting from our clean-trained checkpoints, we fine-tune with reduced learning rates: CNNs use SGD with LR 10^{-3} and cosine annealing for up to 100 epochs, while ViTs use AdamW with the same learning rate and schedule. Gradients are clipped to a global L_2 norm of 10, and batches with non-finite loss are skipped. Data augmentation includes random cropping with 4-pixel padding and horizontal flipping; all models consume raw pixel inputs in $[0, 1]$ without mean/std normalization, and attacks and epsilon budgets operate directly in this pixel space. The training batch size is 128 for the CNN and 64 for the ViT (evaluation batch sizes vary by analysis). Model selection and early stopping use PGD-10 adversarial accuracy on a held-out 2,000-sample validation split carved from the training set, so the test set is touched only by the final post-hoc evaluations. The training seeds are 42 for ResNet-18 and 123 for ViT-Tiny; since the validation permutation is derived from the training seed, the two models use different (fixed) held-out splits. Because the validation images were seen during clean pretraining, validation accuracies are optimistic in absolute terms, but this does not affect test-set results. Early stopping uses patience 20 with a 0.1 percentage-point improvement threshold: the ResNet-18 run stopped at epoch 84 and the ViT-Tiny run at epoch 51, truncating their cosine schedules.

F. Evaluation Metrics

a) *Robustness Metrics*: We evaluate models using three standard metrics: clean accuracy (classification performance on unperturbed test images), robust accuracy (classification performance on adversarial examples), and attack success rate (fraction of correctly classified clean images that become misclassified after attack). For white-box evaluations we additionally report the *conditioned fooling rate*—the fraction of clean-correct samples that the attack flips—and its complement, conditioned survival. Because both PGD and AutoAttack count clean-misclassified samples as non-robust in our per-sample logs, absolute robust accuracy factorizes exactly as $\text{robust} = \text{clean} \times \text{conditioned survival}$, which lets us decompose robustness differences into a clean-accuracy component and a conditional-susceptibility component. Finally, we report accuracies on the *both-correct* subset (samples that both compared models classify correctly on clean inputs), which yields an exactly paired comparison on a shared sample set.

b) *Transfer Attack Metrics*: Our primary transfer metric is the *conditioned fooling rate*: among the samples that the target model classifies correctly on clean inputs, the fraction that becomes misclassified under adversarial examples crafted on the source model. Given adversarial examples $\{x_{adv}^S\}$ generated against source model f_S , the conditioned transfer rate to target model f_T is:

$$T_c(f_S \rightarrow f_T) = \frac{\sum_{i=1}^n \mathbf{1}[f_T(x_i) = y_i] \mathbf{1}[f_T(x_{adv,i}^S) \neq y_i]}{\sum_{i=1}^n \mathbf{1}[f_T(x_i) = y_i]} \quad (7)$$

Conditioning transfer evaluation on correctly classified samples follows established practice [26], [35]–[38] and is essential in cross-architecture comparisons: the raw misclassification rate $\frac{1}{n} \sum_i \mathbf{1}[f_T(x_{adv,i}^S) \neq y_i]$ includes the target’s clean

error, so a target with lower clean accuracy appears more “vulnerable” to transfer even when no additional samples are flipped by the attack. We report the raw rate as a secondary metric for transparency and to quantify the size of this confound. Because the conditioning set itself is a protocol choice, we additionally report two established alternatives: the *both-correct* protocol, which conditions on samples that both the source and the target classify correctly [26], [38], and the *successful-source* protocol, which conditions on target-correct samples whose attack succeeded on the source in the white-box sense. As Section IV-B shows, these protocols can yield qualitatively different asymmetry conclusions, so we treat the protocol as part of the reported result rather than an implementation detail. Transfer rates are computed on the full test set (10,000 samples) with per-sample outcomes logged, and we report 95% confidence intervals obtained by percentile bootstrap over samples (10,000 resamples).

c) *Gradient Analysis Metrics*: We analyze gradient properties through three complementary metric families, computing gradients of per-sample losses so that magnitudes are independent of the evaluation batch size. Gradient norm ($\|\nabla_x \mathcal{L}\|_2$ and $\|\nabla_x \mathcal{L}\|_\infty$) measures overall sensitivity magnitude. Gradient sparsity is quantified with scale-invariant measures—Hoyer sparsity, the Gini coefficient of $|g|$, and the fraction of components below 1% of the per-sample maximum—since fixed absolute thresholds conflate gradient scale with structure; Hoyer and Gini satisfy the scale-invariance axioms formalized by Hurley and Rickard [55]. Gradient alignment measures the average pairwise absolute cosine similarity across samples (computed over all pairs of the 500 analysis samples):

$$\text{Alignment} = \frac{1}{N(N-1)} \sum_{i \neq j} \left| \frac{\nabla_{x_i} \mathcal{L} \cdot \nabla_{x_j} \mathcal{L}}{\|\nabla_{x_i} \mathcal{L}\| \|\nabla_{x_j} \mathcal{L}\|} \right| \quad (8)$$

We use the absolute cosine because shared sensitivity *axes* matter for attack geometry regardless of sign; since sign coherence is what universal-perturbation arguments additionally require, we also report the signed mean cosine, which turns out to be near zero for every model (Section IV-C). High absolute alignment indicates similar perturbation axes across inputs.

d) *Feature Degradation Metrics*: To analyze how adversarial perturbations affect ViT representations across layers, we compute three metrics at each layer ℓ . Feature cosine similarity measures directional preservation of the representation:

$$\text{Sim}_\ell = \frac{h_\ell(x) \cdot h_\ell(x_{adv})}{\|h_\ell(x)\| \|h_\ell(x_{adv})\|} \quad (9)$$

where h_ℓ denotes the representation at layer ℓ . Feature norm change captures magnitude effects:

$$\Delta_\ell = \frac{\|h_\ell(x_{adv})\| - \|h_\ell(x)\|}{\|h_\ell(x)\|} \times 100\% \quad (10)$$

Feature L_2 distance ($\|h_\ell(x) - h_\ell(x_{adv})\|_2$) quantifies absolute deviation.

G. Statistical Validation

To assess evaluation stochasticity, we evaluate each trained model 3 times with different random seeds (42, 123, 456),

TABLE I: Main Robustness Results on CIFAR-10 ($\epsilon = 8/255$). All cells for our models are measured post hoc on the full 10,000-sample test set with fully seeded evaluations (seed 42, deterministic CUDA). Checkpoints were selected on a held-out validation split (Section III-E). AA = AutoAttack (standard suite).

Model	Defense	Params	Clean	FGSM	PGD-10	AA
<i>CNN Models</i>						
ResNet-18	None	11.2M	94.37	30.93	0.01	–
ResNet-18	AT	11.2M	85.42	49.90	40.97	35.74
WRN-28-10	AT [†]	36.5M	89.48	70.91	66.92	62.76
<i>Vision Transformer Models</i>						
ViT-Tiny	None	5.7M	77.50	3.34	0.03	–
ViT-Tiny	AT	5.7M	75.65	40.07	36.09	32.94

[†]Robust checkpoint of Gowal et al. [48] via RobustBench [17].

Clean/FGSM/PGD-10: our local full-test evaluations; AA: the RobustBench-reported figure (we did not rerun AutoAttack on it).

reporting mean $\pm$ standard deviation. The evaluated subset and its ordering are fixed (the test loader does not shuffle), so the seeds vary *only* the random initialization of PGD restarts; the measured variance therefore quantifies attack-initialization stochasticity, not data-sampling or training variance. To address training-level reproducibility, we conducted multiple independent adversarial-training runs per architecture: two under a preliminary implementation of the training loop, and one final run per architecture under the standardized protocol reported in this paper—rankings were preserved across all runs, with non-negligible run-to-run variation for the ViT (Section IV-E). All evaluation and analysis scripts fix all random seeds (seed 42, including CUDA determinism flags), so every reported number is exactly reproducible from the released pipeline; the 3-seed analysis above bounds the residual attack-initialization stochasticity that an unseeded rerun would exhibit at roughly ± 0.1 percentage points. All experiments are conducted on NVIDIA RTX 5060 Ti (16GB) using PyTorch 2.6.0 with CUDA 12.8.

IV. EXPERIMENTS

A. Robustness Comparison

Table I presents the main robustness comparison across all evaluated models. We report clean accuracy and robust accuracy under PGD-10 and AutoAttack with $\epsilon = 8/255$.

The results indicate a statistically significant *absolute* robustness advantage for the CNN under standard adversarial training: on the full test set, ResNet-18 AT achieves 35.74% robust accuracy under AutoAttack, compared to 32.94% for ViT-Tiny AT—a 2.80 percentage-point gap. Because both models are evaluated on the same 10,000 samples with per-sample outcomes logged, we can apply McNemar’s paired test: 940 samples are robustly classified only by the CNN versus 660 only by the ViT ($\chi^2 = 48.65$, $p < 10^{-11}$), so the gap is not attributable to evaluation noise. The 95% binomial confidence half-widths are below ± 1 percentage point per model. Two caveats remain. First, the gap is modest relative to training-level variability: across our training runs the ViT’s PGD-10 accuracy varied by roughly four points (Section IV-E),

TABLE II: Conditional decomposition of white-box robustness on the full test set. Cond. fooling = fraction of clean-correct samples flipped by the attack (95% bootstrap CIs, 10,000 resamples); absolute robust accuracy factorizes exactly as clean $\times$ (1 – cond. fooling). Both-correct columns evaluate the $n=7,260$ samples that *both* models classify correctly on clean inputs.

Model	Cond. fooling (%)		Both-correct robust (%)	
	PGD-10 [§]	AA	PGD-10 [§]	AA
ResNet-18 AT	52.15 [51.1, 53.2]	58.16 [57.1, 59.2]	54.92	48.15
ViT-Tiny AT	52.33 [51.2, 53.5]	56.46 [55.4, 57.6]	49.61	45.33

[§]PGD-10 columns derive from the per-sample-logged evaluation run (full-test 40.87%/36.06%); Table I reports a separately seeded rerun that agrees within ± 0.1 points.

so the architecture comparison should be read as “under a matched adversarial objective and threat model” rather than as a universal constant. Second, the CNN also retains a higher clean accuracy (85.42% vs 75.65%, a 9.8-point gap), reflecting the difficulty of training ViTs on small datasets from scratch. AutoAttack robustness is consistently lower than PGD-10 robustness for both architectures (5.23 points for ResNet-18, 3.15 points for ViT-Tiny), confirming the importance of using this stronger benchmark. The WideResNet-28-10 baseline [48], at 62.76% robust accuracy, demonstrates how far capacity, additional training data, and a stronger training recipe jointly push robustness beyond architectural choice alone.

Table II decomposes these absolute scores into their clean-accuracy and conditional-susceptibility components (Section III-F). The decomposition changes the reading of the headline result in three ways. Under PGD-10, the conditioned fooling rates of the two models are statistically indistinguishable (52.15% vs 52.33%, overlapping CIs), so the entire absolute PGD gap is accounted for by the clean-accuracy gap: 85.42×0.478 vs 75.65×0.477 . Under AutoAttack, the per-model conditional ordering even reverses slightly in the ViT’s favor (58.16% vs 56.46% fooling, CIs nearly touching): on the samples it solves cleanly, the ViT loses proportionally *less* to the stronger attack ensemble. The jointly solved subset tells a third story: restricted to the 7,260 samples both models classify correctly—an exactly paired comparison on shared inputs—the CNN retains a significant advantage under both attacks (PGD-10: 54.92% vs 49.61%, McNemar $\chi^2 = 88.1$, $p < 10^{-20}$; AutoAttack: 48.15% vs 45.33%, $\chi^2 = 27.4$, $p < 10^{-6}$). These views reconcile compositionally: the samples that only the CNN solves cleanly ($n=1,282$) are almost never robust (below 8% PGD survival), which dilutes the CNN’s per-model conditional rate down to the ViT’s level. The headline “the CNN is more robust” therefore holds absolutely and on the common subset, while per-model conditional susceptibility is essentially matched under PGD-10 and slightly ViT-favorable under AutoAttack—a multi-view reading we carry through the rest of the paper.

Figure 1 visualizes this robustness gap across architectures. To further understand the effect of perturbation magnitude, Figure 2 shows PGD-10 robust accuracy across perturbation budgets $\epsilon \in \{2/255, 4/255, 8/255, 16/255\}$ on the full test set.




Fig. 1: Clean, FGSM, PGD-10, and AutoAttack accuracy at $\epsilon = 8/255$ on the full CIFAR-10 test set. CNNs outperform ViTs under standard adversarial training; WideResNet-28-10 achieves the highest robust accuracy through capacity, additional training data, and a stronger recipe jointly [48] (its AutoAttack bar is the RobustBench-reported value).




Fig. 2: PGD-10 robust accuracy vs. perturbation budget (ϵ), measured on the full test set. Both AT architectures degrade with similar shapes as epsilon increases, with the CNN maintaining a consistent margin.

Figure 3 provides visual examples of adversarial perturbations and their effects on model predictions.

B. Transfer Attack Analysis

We analyze the transferability of adversarial examples between architectures using PGD-10 attacks on the full test set ($n=10,000$). Table III reports the conditioned fooling rate defined in Section III-F—the fraction of samples the target classifies correctly on clean inputs that become misclassified under adversarial examples crafted on the source—with raw misclassification rates in parentheses for transparency.

The multi-protocol analysis (Table IV) yields a sharper methodological finding than either single metric would. The raw rates show an apparent 8.3-point asymmetry favoring CNN sources (40.2% vs 31.9%): because the ViT target has roughly ten points more clean error than the CNN target, the raw metric credits the CNN source with failures the ViT would have made on clean inputs anyway. Under the target-



Fig. 3: Adversarial example visualization. Each row shows: original image with model prediction, perturbation magnified $10\times$, and adversarial image with new prediction. Row 1: Attack on ResNet-18. Row 2: Attack on ViT-Tiny. Row 3: Transfer attack (generated on ResNet, evaluated on ViT). The perturbations are bounded by $\epsilon=8/255$ and visually subtle at the displayed scale, yet cause confident misclassifications.

TABLE III: Transfer Attack Matrix on the full test set ($n=10,000$). Cells show the conditioned fooling rate with 95% bootstrap confidence intervals (10,000 resamples); raw target-misclassification rates in parentheses. Diagonal entries are white-box attacks.

Source $\rightarrow$ Target	ResNet-18 AT	ViT-Tiny AT
ResNet-18 AT	52.2 [51.1, 53.3] (59.1)	21.0 [20.0, 21.9] (40.2)
ViT-Tiny AT	20.3 [19.5, 21.2] (31.9)	52.3 [51.3, 53.4] (63.9)

TABLE IV: Transfer asymmetry under three conditioning protocols (PGD-10, $\epsilon=8/255$, full test set; 95% bootstrap CIs, 10,000 resamples). Both-correct paired analysis: difference 5.33 points, paired bootstrap CI [4.39, 6.25], sign-flip permutation $p < 5 \times 10^{-5}$; TOST equivalence fails even at a ± 3 -point margin.

Protocol	CNN $\rightarrow$ ViT	ViT $\rightarrow$ CNN	Diff.
Target-correct	20.95 [20.03, 21.88]	20.32 [19.47, 21.19]	+0.63
Both-correct	19.19 [18.26, 20.10]	13.86 [13.07, 14.66]	+5.33
Successful-source	39.07 [37.47, 40.66]	33.79 [32.48, 35.13]	+5.28

correct conditioned metric this asymmetry vanishes: 20.95% vs 20.32%, a 0.63-point difference that is not significant ($z = 0.98$, $p = 0.33$) and is formally equivalent within a ± 2 -point margin (TOST, $p = 0.016$; equivalence also holds at ± 3 but not at ± 1 , $p = 0.72$)—and this target-correct symmetry is stable under the momentum-based MI-




Fig. 4: Conditioned (target-correct) transfer fooling-rate heatmap ($n=10,000$). Off-diagonal rates are near-parity under this protocol (21.0% vs 20.3%); Table IV shows that stricter protocols reveal a directional difference.

FGSM [36] (19.85% [18.97, 20.78] vs 20.11% [19.28, 20.96]). However, the two protocols that restrict the comparison further *reverse the conclusion again*: on the both-correct subset the directions separate decisively (19.19% vs 13.86%, a 5.33-point CNN $\rightarrow$ ViT advantage; paired bootstrap CI [4.39, 6.25], sign-flip permutation $p < 5 \times 10^{-5}$, TOST equivalence rejected even at ± 3 points), and the successful-source protocol agrees (5.28-point gap). The mechanism is compositional: samples that the target solves cleanly but the *source* does not are hyper-vulnerable in both directions (56.9% fooling for ViT $\rightarrow$ CNN on its $n=1,282$ such samples; 63.0% for CNN $\rightarrow$ ViT on its $n=305$), and the asymmetric sizes of these marginal subsets happen to pull the two target-correct rates to parity. Applying the identical target-correct protocol to the *clean-trained* counterparts underlines the same lesson from another angle: conditioned transfer is much lower (9.91% [9.32, 10.53] CNN $\rightarrow$ ViT vs 11.76% [11.15, 12.41] ViT $\rightarrow$ CNN, a small but distinguishable asymmetry), while the raw metric would report a spurious 12.2-point gap in the *opposite* direction (28.3% vs 16.1%). Taken together, the cautionary result is stronger than a clean-error confound alone: *the measured direction and size of cross-architecture transfer asymmetry depend on the conditioning protocol*, which can create (raw), erase (target-correct), or reveal (both-correct, successful-source) a directional difference on the same underlying data. Cross-architecture transfer studies should therefore state their protocol explicitly and, where feasible, report more than one; none of the protocol readings here licenses interpreting raw-rate asymmetries architecturally.

Figure 4 visualizes the conditioned transfer matrix.

C. Gradient Characteristics

To characterize architectural differences in input sensitivity, we analyze the gradient properties of both architectures, together with a CIFAR-native ViT control (32×32 inputs, 4×4

TABLE V: Gradient Characteristics ($n=500$ test samples; per-sample loss gradients; alignment is the mean absolute cosine similarity over *all* sample pairs—random 50-sample partition means agree within 0.0002; $\pm$ values are std across 10 batch means). All ResNet–ViT differences on the scale-invariant sparsity metrics are significant under per-sample paired Wilcoxon tests (Holm-corrected $p < 10^{-15}$; paired Cohen’s $d = 0.34$ – 0.65).

Metric	ResNet-18 AT	ViT-Tiny AT	ViT (native) [‡]
L_2 Norm (mean)	0.698	0.760	0.554
L_∞ Norm (mean)	0.139	0.144	0.087
Hoyer Sparsity	0.474 ± 0.009	0.449 ± 0.011	0.405 ± 0.008
Gini Coefficient	0.634 ± 0.006	0.606 ± 0.008	0.571 ± 0.006
Frac. $< 1\%$ of max	31.5% ± 1.7	25.8% ± 1.7	17.9% ± 1.2
Gradient Alignment	0.038 ± 0.002	0.052 ± 0.002	0.079 ± 0.005

[‡]CIFAR-native ViT control (4×4 patches, no upsampling; 5.4M params; checkpoint-selection clean/robust acc. 62.4%/31.8%). Included only to test the upsampling confound.

patches, no upsampling) trained with the same adversarial protocol, which isolates the effect of the $32 \rightarrow 224$ upsampling in our main ViT pipeline. Table V summarizes the key statistics using scale-invariant sparsity measures (Section III-F); gradients are computed with per-sample losses so that norms do not depend on the evaluation batch size.

The gradient analysis reveals consistent but modest architectural differences. Across all three scale-invariant sparsity measures, CNN gradients are sparser than ViT gradients (Hoyer 0.474 vs 0.449; Gini 0.634 vs 0.606; 31.5% vs 25.8% of components below 1% of the per-sample maximum), indicating that CNN sensitivity is somewhat more concentrated in fewer input components. Because the same 500 samples are measured under both models, we test these differences with per-sample paired statistics: all three are highly significant (Wilcoxon signed-rank, Holm-corrected $p < 10^{-15}$) with small-to-medium paired effect sizes (Cohen’s d of 0.34, 0.65, and 0.43 for Hoyer, Gini, and the relative-threshold fraction), quantifying “consistent but modest” precisely. Interestingly, this ordering is a product of adversarial training: in the clean-trained counterparts the ViT is the sparser model (Hoyer 0.406 vs 0.354), and adversarial training sparsifies the CNN far more strongly (+0.12 Hoyer) than the ViT (+0.04), consistent with the known sparsifying effect of adversarial training on input gradients [22]. The native-resolution ViT control is *less* sparse than the upsampled ViT on every measure, so the CNN–ViT sparsity ordering is not an artifact of the $32 \rightarrow 224$ upsampling—if anything, upsampling moves the ViT *toward* the CNN. We emphasize that these differences are far smaller than a threshold-based sparsity metric would suggest: an earlier version of this analysis (fixed 10^{-6} threshold, measured on the preliminary run2 models) indicated a $1.7\times$ sparsity ratio, most of which reflected gradient-scale differences rather than structure; on the final models the same scale-dependent metric inflates the ratio even further ($\approx 2.4\times$). Cross-sample gradient alignment is low for both architectures but measurably higher for the ViT (0.052 vs 0.038 in mean absolute cosine over all pairs, a $1.36\times$ ratio; native control 0.079), suggesting that transformer sensitivity axes are somewhat more shared




Fig. 5: Gradient characteristics comparison. (a) Input-gradient magnitude maps for representative test images with the corresponding PGD-10 perturbations. (b) Distribution of per-sample gradient L_2 norms for both architectures.




Fig. 6: Gradient heatmap comparison between CNN and ViT on the same input images. CNN gradients (top) are modestly more concentrated around object edges, while ViT gradients (bottom) are somewhat more spatially distributed (Hoyer sparsity annotations from the gradient-analysis artifacts).

across inputs. The *signed* mean cosine, by contrast, is near zero for every model (≤ 0.002): what is shared across samples is sensitivity axes, not consistent directions, so the alignment gap should not be read as direct evidence of cheap universal perturbations. We treat both characteristics as behavioral correlates of, rather than established explanations for, the architectures’ differing adversarial behavior.

Figure 5 provides visual evidence for these gradient differences.

Figure 6 shows a direct comparison of gradient heatmaps, highlighting the structural differences in how each architecture computes input gradients.

D. Feature Degradation in Vision Transformers

We analyze how adversarial perturbations affect ViT feature representations across layers, extending layer-wise representation analyses of clean-trained ViTs [23] to the adversarially trained setting. The measured representations are the outputs of each transformer block’s MLP sub-block; the attention-weight

TABLE VI: ViT-Tiny Feature Degradation Under PGD-10 Attack ($n=100$ test samples; features are MLP sub-block outputs; $\pm$ values are std across 5 batch means)

Layer	L_2 Distance	Cosine Sim.	Norm Change
Block 0	2.73 ± 0.13	0.994 ± 0.000	$-0.11 \pm 0.17\%$
Block 1	1.74 ± 0.07	0.992 ± 0.001	$-0.05 \pm 0.22\%$
Block 2	1.81 ± 0.05	0.981 ± 0.002	$-0.99 \pm 0.31\%$
Block 3	1.33 ± 0.05	0.974 ± 0.002	$-0.69 \pm 1.24\%$
Block 4	1.71 ± 0.08	0.972 ± 0.002	$-1.07 \pm 1.76\%$
Block 5	1.92 ± 0.03	0.972 ± 0.003	$-0.29 \pm 1.54\%$
Block 6	1.84 ± 0.04	0.961 ± 0.004	$-0.28 \pm 0.83\%$
Block 7	1.48 ± 0.06	0.963 ± 0.004	$-1.40 \pm 0.99\%$
Block 8	1.79 ± 0.03	0.955 ± 0.003	$-1.05 \pm 0.27\%$
Block 9	1.86 ± 0.02	0.963 ± 0.002	$-0.73 \pm 0.54\%$
Block 10	1.73 ± 0.09	0.965 ± 0.003	$-0.76 \pm 0.45\%$
Block 11	2.31 ± 0.15	0.965 ± 0.004	$-1.11 \pm 0.80\%$

analysis accompanies the figure-level results. Table VI shows all twelve blocks ($n=100$ test samples).

The layer-wise analysis reveals a characteristic degradation profile. Cosine similarity between clean and adversarial features drops quickly in the early blocks (0.994 at Block 0 to 0.972 by Block 4), reaches its minimum in the middle-to-late network (0.955 at Block 8), and then plateaus with a slight recovery in the final blocks (0.965 at Blocks 10–11). Most of the directional deviation is therefore accumulated by mid-network rather than growing monotonically to the output—an “early-drop, mid-network-plateau” profile rather than strictly progressive degradation. The final-block recovery is statistically reliable at the batch level (Block 8 vs Block 11 cosine similarity, Welch $t = -4.07$, $p = 0.005$). A parallel analysis of the adversarially trained ResNet-18’s eight residual-block outputs (same attack, same metrics, $n=100$) shows a qualitatively different profile: cosine similarity degrades *monotonically* with depth (0.995 at layer1.0 down to 0.913 at layer4.1, with no plateau or recovery), accompanied by much larger norm distortions in the final blocks (-6.4% at layer4.0, $+7.2\%$ at layer4.1). The plateau-and-recovery shape with preserved norms is therefore specific to the evaluated ViT in our model pair, not a generic property of deep networks under attack. Feature norm changes remain modest throughout (at most -1.4%), indicating that adversarial perturbations corrupt the *direction* of representations far more than their magnitude. One untested hypothesis is that the mild recovery in the final blocks reflects a regularizing influence of the classifier head on the last transformer blocks. Even so, the terminal deviation (cosine similarity ≈ 0.965) is sufficient to flip predictions on most attacked samples: on these 100 samples, accuracy drops from 79% clean to 39% under PGD-10.

Complementing the feature-level measurements, Figure 7 visualizes CLS-token attention maps at layers 1, 6, and 12 for a representative sample: early-layer attention is nearly unchanged under attack, while later layers show visibly larger attention shifts, consistent with attention-shift observations under patch perturbations [28] and with attacks that explicitly target the attention mechanism [27].

Figure 8 complements the feature-level analysis with attention-level measurements. Notably, the per-layer *entropy* of the CLS attention distribution is nearly unchanged under




Fig. 7: CLS-token attention maps (post-softmax, head-averaged) across ViT layers under adversarial attack. Each row shows a different layer (1, 6, 12); columns show clean attention, adversarial attention, and their difference. Early layers maintain similar attention patterns, while later layers show larger attention shifts.




Fig. 8: ViT attention analysis under attack: (a) clean vs. adversarial CLS attention for a representative sample; (b) per-layer entropy of the CLS attention distribution for clean and adversarial inputs.

attack (mean absolute change ≈ 0.015 nats on layer means of ≈ 4.9 nats, at most 0.8%; measured on the 8-sample figure batch): the attack redirects *where* attention goes (Figure 7) without broadening or sharpening its distribution, reinforcing that the operative degradation is directional drift in the feature space rather than attention diffusion.

Figure 9 shows an illustrative t-SNE embedding of the penultimate feature spaces; because t-SNE embeddings are sensitive to hyperparameters, we base all claims on quantitative metrics computed in the original feature space ($n=500$, white-box PGD-10 per model). These metrics show that the attack collapses class structure most strongly in the *CNN*: clean-fit 5-NN label consistency drops from 85.2% to 47.8%,




Fig. 9: Illustrative t-SNE embedding of penultimate features for clean samples (circles) and white-box PGD-10 adversarial samples (triangles), for ResNet-18 (left) and ViT-Tiny (right). t-SNE is shown for orientation only; the quantitative comparison is given in the text.

TABLE VII: Statistical Validation (Mean $\pm$ Std over 3 evaluation seeds, first 1,000 test samples)

Model	Clean Acc (%)	PGD-10 Acc (%)
ResNet-18 AT	86.10 $\pm$ 0.00	40.57 $\pm$ 0.12
ViT-Tiny AT	76.90 $\pm$ 0.00	38.23 $\pm$ 0.09

the silhouette score falls from 0.14 to -0.06 , the intra/inter-class distance ratio rises from 0.64 to 0.91, and the relative class-centroid shift is 0.29. The ViT’s penultimate space, by contrast, is only weakly class-separated even on clean inputs (silhouette ≈ 0 ; 5-NN consistency 68.2%, dropping to 53.4% under attack) and its centroids shift far less (0.09). An earlier qualitative reading of these plots (“more class-coherent CNN displacement, greater ViT inter-class confusion”) is not supported by the quantitative metrics and we do not maintain it; the measurable statement is that the CNN’s cleanly well-separated feature space is heavily disrupted by white-box attack, whereas the ViT’s feature-space class structure is weaker to begin with and degrades less in these aggregate terms.

E. Statistical Validation

To quantify evaluation stochasticity, we evaluate each trained model across 3 different random seeds (42, 123, 456) for attack initialization on 1,000 test samples. Table VII shows the results.

The low variance confirms that evaluation stochasticity is negligible. The evaluated subset and its ordering are identical across seeds (the test loader does not shuffle), so clean accuracy shows exactly zero variance and the only stochastic element is the random initialization of PGD restarts. Robust accuracy shows minimal variation (standard deviation 0.12 percentage points for CNN, 0.09 for ViT), indicating stable attack behavior across random restarts. This procedure quantifies attack-initialization stochasticity only; it does not measure data-sampling or training-level variance. (Differences between Table VII and Table I PGD values—0.4 points for the CNN and up to 2.1 points for the ViT—reflect the easier first-1,000-sample subset versus the full test set.)

Beyond evaluation seeds, we conducted multiple independent adversarial-training runs per architecture during this study. Under the final training protocol, ResNet-18 AT reached

40.97% and ViT-Tiny AT 36.09% full-test PGD-10 accuracy; earlier runs under a preliminary implementation of the training loop (differing in attack-generation details, later standardized as described in Section III-D) produced values in the 40.3–41.0% (CNN) and 32.8–36.9% (ViT) ranges. The CNN-over-ViT ranking was preserved in every run, though the ViT’s run-to-run variation is clearly larger than evaluation-level variance, and we therefore avoid attributing precise magnitudes to the architecture gap beyond the paired analysis of Section IV-A.

Computational Details: All experiments were conducted on an NVIDIA RTX 5060 Ti (16GB VRAM) using PyTorch 2.6.0 with CUDA 12.8. AutoAttack evaluation used the standard configuration (APGD-CE, APGD-T, FAB-T, Square) on the full 10,000-sample test set, evaluated in seeded 1,000-sample chunks with per-sample outcome logging; total wall-clock time was several hours per model on this hardware.

V. DISCUSSION

A. Key Findings and Their Implications

a) Transfer Asymmetry Is a Property of the Measurement Protocol as Much as of the Models: Our experiments sharpen a methodological point that we believe deserves more attention in cross-architecture transfer studies. Conditioning transfer evaluation on correctly classified samples is established practice [35], [36], [38]; what our controlled same-experiment comparison adds is the demonstration that *which* conditioning protocol one adopts can qualitatively change the conclusion on the same data (Table IV). The raw metric manufactures an 8.3-point asymmetry favoring CNN sources out of the targets’ ten-point clean-accuracy difference. The target-correct conditioned metric removes that confound and shows near-parity (20.95% vs 20.32%, TOST-equivalent within ± 2 points). Yet the both-correct and successful-source protocols—the former used by careful prior work [26], [38]—reveal a significant 5.3-point CNN $\rightarrow$ ViT advantage on the jointly solved samples. No single number here is “the” transfer asymmetry: the raw gap is an artifact, the target-correct parity reflects a compositional cancellation (marginal samples that only the target solves are hyper-vulnerable in both directions and asymmetrically numerous), and the both-correct gap is the cleanest paired estimate of directional transferability on shared inputs. We therefore recommend that cross-architecture transfer studies state their conditioning protocol explicitly, justify it, and where feasible report several—treating protocol sensitivity itself as a result rather than a nuisance.

This reframing also changes the interpretive link to our gradient measurements. Under the both-correct protocol, adversarial examples crafted on the sparser-gradient CNN transfer better to the ViT than vice versa; the direction co-occurs with the gradient-structure differences we measure, but our correlational design cannot attribute it, and the target-correct parity shows the aggregate picture is protocol-dependent. Mahmood et al. [26] observe cross-architecture asymmetries under a both-correct protocol for normally trained models; our both-correct result is directionally consistent with theirs, while our raw-metric analysis shows how easily a differently manufactured asymmetry could be mistaken for such a finding when targets differ in clean accuracy.

b) Gradient Characteristics as Vulnerability Signatures:

The gradient analysis characterizes how the two architectures distribute input sensitivity. Across three scale-invariant measures [55], CNN gradients are consistently sparser than ViT gradients (Hoyer 0.474 vs 0.449; Gini 0.634 vs 0.606; 31.5% vs 25.8% of components below 1% of the per-sample maximum; paired Wilcoxon Holm-corrected $p < 10^{-15}$, paired Cohen’s d 0.34–0.65), indicating modestly more concentrated sensitivity—consistent with the known tendency of adversarial training to sparsify input gradients [22] and to reshape gradient structure more broadly [21], here differing in degree between architectures. We deliberately phrase this as concentration in fewer input *components*: whether the concentration is also spatially localized is a distinct claim that our sparsity measures do not test. Two controls strengthen this reading. First, the native-resolution ViT control is even less sparse than the upsampled ViT, so the ordering is not an artifact of input upsampling. Second, the effect size is honest: an earlier threshold-based analysis (on the preliminary run2 models) suggested a $1.7\times$ sparsity ratio, most of which we traced to gradient-scale differences rather than structure—the structural difference is real but small. In parallel, cross-sample gradient alignment is measurably higher for the ViT (0.052 vs 0.038, all-pairs absolute cosine), suggesting somewhat more shared sensitivity directions across inputs, though both values are low in absolute terms.

The alignment result merits careful interpretation. Universal adversarial perturbations [56] exploit sensitivity directions shared across inputs; one might therefore hypothesize that ViTs’ global self-attention produces higher cross-sample gradient alignment. We observe a difference in that direction—ViT alignment is $1.36\times$ the CNN’s (0.052 vs 0.038), and the native-resolution ViT control is higher still (0.079)—but two caveats bound the interpretation. First, the absolute levels are low for both architectures, far from the strong alignment that would make universal perturbations cheap. Second, the *signed* mean cosine is near zero for every model (≤ 0.002): samples share sensitivity axes but not consistent directions, whereas universal perturbations require sign-coherent alignment—so the measured gap concerns the geometry of sensitivity subspaces rather than directly predicting universal-perturbation vulnerability. Measuring the same quantity in the clean-trained counterparts resolves the training-regime question directly: clean alignment is 0.050 for the ViT and 0.030 for the CNN, essentially the same levels as after adversarial training. The ViT–CNN alignment gap is therefore an *architectural* property that adversarial training neither creates nor suppresses appreciably in our setting—in contrast to gradient magnitude, which adversarial training reduces by an order of magnitude for both models, and to sparsity, whose architecture ordering adversarial training reverses (Section IV-C).

c) Feature Degradation in Vision Transformers: Our layer-wise analysis reveals a characteristic feature degradation profile in ViT layers. In the first two transformer blocks, adversarial inputs produce noticeable L_2 distance from clean representations while maintaining semantic direction (cosine similarity above 0.99) and near-constant norms; the perturbation initially displaces representations without appreciably

altering either their norm or their semantic direction. Directional deviation then accumulates rapidly through the early-to-middle network, with cosine similarity reaching its minimum of 0.955 at Block 8, after which it plateaus with a mild recovery in the final blocks (0.965 at Blocks 10–11); one untested hypothesis for the recovery is a regularizing influence of the classifier head. The degradation is subtler than the attack’s effectiveness might suggest—norm changes remain below 1.5% throughout—implying that modest directional shifts, largely established by mid-network, are sufficient to flip predictions on most attacked samples.

This early-drop, mid-network-plateau profile suggests that effective defenses should focus on stabilizing early-to-middle layer representations rather than only filtering perturbations at the input level or protecting the final layers. Intermediate supervision or layer-wise consistency regularization might help maintain representational integrity throughout the network. The parallel residual-block analysis of the adversarially trained CNN (Section IV-D) confirms the profile is specific to the evaluated model pair rather than generic: the CNN’s representations degrade monotonically with depth and with large terminal norm distortions, suggesting the two models would benefit from differently placed stabilization.

B. Broader Implications

a) For Researchers: Our work demonstrates the value of multi-view behavioral analysis over single-number benchmarking. Simply reporting that one architecture achieves higher robust accuracy than another provides limited insight: as our conditional decomposition shows, the same models support the readings “the CNN is more robust” (absolute and both-correct views), “conditional susceptibility is matched” (per-model PGD view), and “the ViT loses less on its solved samples” (per-model AutoAttack view), depending on the conditioning. Characterizing *how* architectures behave differently—through protocol-explicit transfer analysis, gradient characterization, and feature degradation tracking—enables targeted improvements rather than trial-and-error defense development.

The consistent gap between AutoAttack and PGD-10 results (3.1–5.2 percentage points lower robust accuracy under AutoAttack for our models; Section IV-A) reinforces the importance of using standardized, strong evaluation protocols. While PGD-10 remains useful during development for its computational efficiency, final robustness claims should always be validated with AutoAttack or equivalent comprehensive evaluation. Beyond this, conditioned transfer rates measured from diverse source architectures could serve as a diagnostic for architecture-specific vulnerabilities—a direction our two-model study cannot itself establish, but one that the clean-accuracy-controlled methodology makes feasible.

b) For Practitioners: For deployment scenarios requiring adversarial robustness, our findings offer concrete guidance. Under matched standard adversarial training on this benchmark, the CNN provides better absolute robust accuracy with more stable training dynamics; note, however, that most of its absolute PGD advantage reflects its higher clean accuracy rather than lower conditional susceptibility, so the practical

choice also depends on which accuracy profile a deployment values. ViTs require more careful hyperparameter tuning during adversarial training and may benefit from specialized techniques beyond standard adversarial training.

When deploying ensemble defenses that combine CNN and ViT models, practitioners should evaluate transfer under an explicitly chosen, clean-accuracy-controlled protocol—and ideally more than one—before assuming either architecture is the harder surrogate; our analysis shows that the raw metric manufactures an asymmetry, the target-correct metric shows parity, and the both-correct metric reveals a genuine CNN→ViT advantage on jointly solved samples. Additionally, the gap between WideResNet-28-10 (62.76% robust accuracy) and ResNet-18 (35.74%) shows how much headroom lies beyond our baselines—though this particular model owes its robustness to capacity, additional (unlabeled) training data, and a stronger training recipe jointly [48], so the gap should not be read as a capacity effect alone. For high-stakes applications, investing in larger models trained with strong recipes and augmented data may prove more effective than developing novel defense mechanisms.

c) On Model Capacity: Our comparison involves models with different parameter counts: ResNet-18 has 11.2M parameters while ViT-Tiny has 5.7M. This $2\times$ capacity difference could partially explain the robustness gap. However, our mechanistic analyses suggest that architectural factors play a more fundamental role than model size alone.

While parameter count differs (11.2M vs. 5.7M), our gradient analysis reveals scale-invariant structural differences—modestly sparser CNN gradients and moderately higher ViT absolute alignment—that are consistent with architectural processing differences and survive the native-resolution control. The transfer evidence is protocol-dependent (parity under target-correct conditioning, a CNN→ViT advantage under both-correct conditioning), so whether gradient-level differences translate into directional transfer effects cannot be settled by our design; architectural attribution of the robust-accuracy gap itself likewise remains partial: capacity, clean-accuracy differences, and training-recipe maturity for ViTs on small datasets all plausibly contribute. We acknowledge that capacity-controlled experiments (e.g., using ViT-Small with ~ 22 M parameters) would provide stronger evidence for disentangling these factors. We leave such experiments for future work.

C. Limitations

Several limitations constrain the scope of our conclusions and suggest directions for future validation.

Our ViT-Tiny model uses 224×224 bilinear upsampling from CIFAR-10’s native 32×32 resolution. Although this input pipeline mirrors how ImageNet-scale ViT architectures are commonly adapted to small images, our model is trained from scratch (no ImageNet pretraining), so the upsampling is an architectural-adaptation choice rather than a requirement of pretrained weights. Upsampling could in principle confound input-gradient statistics, since each input pixel influences many backbone positions. Our CIFAR-native ViT control

(Section IV-C) addresses this directly: the native-resolution model is even less sparse and more aligned than the upsampled one, so the CNN–ViT ordering of gradient characteristics is not an upsampling artifact. The control model differs in capacity and patch size, however, so it calibrates the preprocessing confound rather than providing a second matched comparison; the semantic drift analysis likewise reflects the upsampled pipeline.

Our evaluation focuses exclusively on CIFAR-10, and findings may differ on ImageNet where ViTs typically excel and pretrained robust models are available. We evaluate only L_∞ attacks; L_2 or semantic attacks may reveal different vulnerability patterns. We validate evaluation stability across three random seeds for attack initialization (standard deviation below 0.25 percentage points) and report multiple training runs per architecture (Section IV-E), though additional training runs under the final protocol would be needed to attach uncertainty estimates to the architecture comparison itself. Finally, standard adversarial training for ViT produces lower robust accuracy than specialized techniques such as diffusion-based data augmentation, with which ViTs have recently been reported to surpass WideResNets on small-scale datasets [31], [57]; our ViT results represent baseline adversarial training performance rather than state-of-the-art ViT robustness, and our architecture comparison holds specifically under matched standard adversarial training.

D. Future Directions

Our findings motivate several research directions. Systematically evaluating CNN–ViT hybrid architectures such as ConViT and CoAtNet could reveal whether combining architectural elements yields robustness benefits beyond either paradigm alone. The gradient sparsity difference we identify—CNN gradients modestly more concentrated under scale-invariant measures, ViT gradients somewhat more spatially distributed—suggests that hybrids might achieve more balanced vulnerability profiles.

Developing training losses that explicitly encourage attention pattern stability under perturbation may address the vulnerability correlates we characterize in ViTs. The early-accumulating semantic drift we observe (largely established by mid-network) suggests that intermediate-layer supervision or consistency regularization could help maintain feature integrity throughout the network.

Investigating whether the robustness gap persists at larger scales (ViT-Base, ViT-Large) and on higher-resolution datasets would establish the generality of our conclusions. Finally, extending our analysis to certified defense methods, including randomized smoothing and interval bound propagation, would determine whether the architectural differences we identify affect achievable certification radii.

VI. CONCLUSION

This paper presented a behavioral characterization of how CNN and Vision Transformer architectures differ under adversarial attack. Moving beyond simple robustness comparisons, we examined these differences through transfer attack analysis,

gradient characterization, and layer-wise feature degradation studies.

Our analysis establishes a statistically significant but modest *absolute* robustness advantage for the CNN under full-test AutoAttack evaluation (35.74% vs 32.94%, McNemar $p < 10^{-11}$), together with a conditional decomposition that qualifies it: the PGD-10 gap is fully accounted for by the clean-accuracy gap (conditioned fooling 52.15% vs 52.33%), the per-model AutoAttack conditional ordering slightly favors the ViT, and yet on the jointly solved sample subset the CNN’s advantage persists under both attacks—so the headline comparison depends on the view, and we report all three. The WideResNet baseline’s much higher robustness further shows that capacity, training data, and training recipe jointly matter beyond architectural choice. A key methodological finding concerns transfer measurement: the apparent cross-architecture asymmetry depends on the conditioning protocol. Raw rates manufacture an 8.3-point asymmetry out of the targets’ unequal clean accuracy; the target-correct conditioned metric shows statistical parity; and the stricter both-correct and successful-source protocols reveal a significant 5.3-point CNN→ViT advantage on jointly solved samples. Transfer studies should therefore state and justify their conditioning protocol explicitly—none of these views licenses interpreting raw-rate asymmetries architecturally. Gradient analysis with scale-invariant measures shows CNN sensitivity is modestly more concentrated—a difference that survives a native-resolution ViT control and per-sample paired tests—while cross-sample gradient alignment is low for both architectures, moderately higher for the ViT in absolute cosine, and near zero in signed cosine for both. We also identify a characteristic degradation profile whereby adversarial perturbations cause directional deviation in ViT feature representations that accumulates by mid-network (cosine similarity falling from 0.994 to 0.955) and then plateaus, despite feature-norm changes below 1.5%—in contrast to the CNN’s monotonic depth profile.

This work characterizes gradient structure alongside cross-architecture transfer behavior, dissociating which gradient properties are architectural (cross-sample alignment) and which are shaped by adversarial training (magnitude and sparsity ordering), and offering methodological guidance in addition to performance comparison. Our systematic layer-wise analysis characterizes how adversarial perturbations propagate through ViT layers, and we distill practical guidance for evaluating and deploying mixed-architecture ensembles. We will release our analysis pipeline to enable future work extending these findings to other architectures and datasets.

Understanding *how* architectures behave differently under adversarial attack—not just *how much*—is essential for developing targeted defenses. Our analysis shows that the behavioral differences between CNNs and ViTs correlate with scale-invariant differences in gradient structure, particularly how sensitivity is distributed across input components—an effect we verified is not an artifact of gradient scale or input preprocessing—though disentangling architectural effects from capacity and training-recipe factors requires the capacity-controlled follow-ups we outline. Cross-sample gradient alignment is low for both adversarially trained models

yet remains moderately higher for the ViT (0.052 vs 0.038, $1.36\times$); measuring the clean-trained counterparts shows nearly identical levels (0.050 vs 0.030), identifying the alignment gap as an architectural property largely independent of the training regime—while gradient sparsity, whose architecture ordering adversarial training reverses, emerges as the training-shaped characteristic. These dissociations open new questions about the interaction between defense methods and architectural inductive biases. We hope this behavioral perspective guides future research toward more principled defense design.

DATA AND CODE AVAILABILITY

The experiments in this study use the CIFAR-10 dataset [46], which is publicly available at <https://www.cs.toronto.edu/~kriz/cifar.html>. Pre-trained WideResNet-28-10 weights were obtained from the RobustBench benchmark [17], [48]. The source code for the analysis pipeline and trained model checkpoints will be made publicly available upon acceptance.

ACKNOWLEDGMENT

The authors acknowledge the computational resources provided by the Department of Software Engineering, Firat University.

REFERENCES

- [1] C. Szegedy, W. Zaremba, I. Sutskever, J. Bruna, D. Erhan, I. Goodfellow, and R. Fergus, “Intriguing properties of neural networks,” in *International Conference on Learning Representations*, 2014.
- [2] I. J. Goodfellow, J. Shlens, and C. Szegedy, “Explaining and harnessing adversarial examples,” in *International Conference on Learning Representations*, 2015.
- [3] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *IEEE Conference on Computer Vision and Pattern Recognition*, 2016, pp. 770–778.
- [4] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [5] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” in *International Conference on Learning Representations*, 2021.
- [6] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, “Towards deep learning models resistant to adversarial attacks,” in *International Conference on Learning Representations*, 2018.
- [7] H. Zhang, Y. Yu, J. Jiao, E. Xing, L. El Ghaoui, and M. Jordan, “Theoretically principled trade-off between robustness and accuracy,” in *International Conference on Machine Learning*, 2019, pp. 7472–7482.
- [8] F. Croce and M. Hein, “Reliable evaluation of adversarial robustness with an ensemble of diverse parameter-free attacks,” in *International Conference on Machine Learning*, 2020, pp. 2206–2216.
- [9] S. Bhojanapalli, A. Chakrabarti, D. Glasner, D. Li, T. Unterthiner, and A. Veit, “Understanding robustness of transformers for image classification,” in *IEEE/CVF International Conference on Computer Vision*, 2021, pp. 10231–10241.
- [10] R. Shao, Z. Shi, J. Yi, P.-Y. Chen, and C.-J. Hsieh, “On the adversarial robustness of vision transformers,” *Transactions on Machine Learning Research*, 2022.
- [11] N. Carlini and D. Wagner, “Towards evaluating the robustness of neural networks,” in *IEEE Symposium on Security and Privacy (SP)*, 2017, pp. 39–57.
- [12] Y. Wang, D. Zou, J. Yi, J. Bailey, X. Ma, and Q. Gu, “Improving adversarial robustness requires revisiting misclassified examples,” in *International Conference on Learning Representations*, 2020.
- [13] C. Guo, M. Rana, M. Cisse, and L. Van Der Maaten, “Countering adversarial images using input transformations,” in *International Conference on Learning Representations*, 2018.

- [14] J. Cohen, E. Rosenfeld, and Z. Kolter, "Certified adversarial robustness via randomized smoothing," in *International Conference on Machine Learning*, 2019, pp. 1310–1320.
- [15] F. Tramèr, A. Kurakin, N. Papernot, I. Goodfellow, D. Boneh, and P. McDaniel, "Ensemble adversarial training: Attacks and defenses," in *International Conference on Learning Representations*, 2018.
- [16] S. Lyu, S. Shaikh, F. Shpilevsky, E. Shelhamer, and M. Lécuyer, "Adaptive randomized smoothing: Certified adversarial robustness for multi-step defences," in *Advances in Neural Information Processing Systems*, vol. 37, 2024.
- [17] F. Croce, M. Andriushchenko, V. Sehwag, E. Debenedetti, N. Flammarion, M. Chiang, P. Mittal, and M. Hein, "RobustBench: a standardized adversarial robustness benchmark," in *NeurIPS Datasets and Benchmarks Track*, 2021.
- [18] D. Wu, Y. Wang, S.-T. Xia, J. Bailey, and X. Ma, "Skip connections matter: On the transferability of adversarial examples generated with ResNets," in *International Conference on Learning Representations (ICLR)*, 2020.
- [19] C. Xie and A. Yuille, "Intriguing properties of adversarial training at scale," in *International Conference on Learning Representations (ICLR)*, 2020.
- [20] C. Xie, M. Tan, B. Gong, J. Wang, A. L. Yuille, and Q. V. Le, "Adversarial examples improve image recognition," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2020, pp. 819–828.
- [21] D. Tsipras, S. Santurkar, L. Engstrom, A. Turner, and A. Madry, "Robustness may be at odds with accuracy," in *International Conference on Learning Representations (ICLR)*, 2019.
- [22] P. Chalasani, J. Chen, A. R. Chowdhury, X. Wu, and S. Jha, "Concise explanations of neural networks using adversarial training," in *Proceedings of the 37th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 119. PMLR, 2020, pp. 1383–1391.
- [23] M. Raghu, T. Unterthiner, S. Kornblith, C. Zhang, and A. Dosovitskiy, "Do vision transformers see like convolutional neural networks?" in *Advances in Neural Information Processing Systems*, vol. 34. Curran Associates, Inc., 2021, pp. 12 116–12 128.
- [24] Y. Bai, J. Mei, A. L. Yuille, and C. Xie, "Are transformers more robust than cnns?" in *Advances in Neural Information Processing Systems*, vol. 34, 2021, pp. 26 831–26 843.
- [25] P. Benz, S. Ham, C. Zhang, A. Karjauv, and I. S. Kweon, "Adversarial robustness comparison of vision transformer and MLP-Mixer to CNNs," in *Proceedings of the 32nd British Machine Vision Conference (BMVC)*, 2021.
- [26] K. Mahmood, R. Mahmood, and M. Van Dijk, "On the robustness of vision transformers to adversarial examples," in *IEEE/CVF International Conference on Computer Vision*, 2021, pp. 7838–7847.
- [27] Y. Fu, S. Zhang, S. Wu, C. Wan, and Y. Lin, "Patch-Fool: Are vision transformers always robust against adversarial perturbations?" in *International Conference on Learning Representations (ICLR)*, 2022.
- [28] J. Gu, V. Tresp, and Y. Qin, "Are vision transformers robust to patch perturbations?" in *Computer Vision – ECCV 2022*, ser. Lecture Notes in Computer Science, vol. 13672. Springer, 2022, pp. 404–421.
- [29] S. Paul and P.-Y. Chen, "Vision transformers are robust learners," in *AAAI Conference on Artificial Intelligence*, 2022, pp. 2071–2081.
- [30] S. Jain and T. Dutta, "Towards understanding and improving adversarial robustness of vision transformers," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2024, pp. 24 736–24 745.
- [31] J. Wu, Z. Song, X. Zhang, S. Xie, L. Lin, and K. Wang, "Vision transformers beat WideResNets on small scale datasets adversarial robustness," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 1, 2025, pp. 886–894.
- [32] Z. Zhu, X. Wang, Z. Jin, J. Zhang, and H. Chen, "Enhancing transferable adversarial attacks on vision transformers through gradient normalization scaling and high-frequency adaptation," in *International Conference on Learning Representations*, 2024.
- [33] N. Papernot, P. McDaniel, and I. Goodfellow, "Transferability in machine learning: from phenomena to black-box attacks using adversarial samples," *arXiv preprint arXiv:1605.07277*, 2016.
- [34] A. Demontis, M. Melis, M. Pintor, M. Jagielski, B. Biggio, A. Oprea, C. Nita-Rotaru, and F. Roli, "Why do adversarial attacks transfer? explaining transferability of evasion and poisoning attacks," in *28th USENIX Security Symposium*, 2019, pp. 321–338.
- [35] Y. Liu, X. Chen, C. Liu, and D. Song, "Delving into transferable adversarial examples and black-box attacks," in *Proceedings of the 5th International Conference on Learning Representations (ICLR)*, 2017.
- [36] Y. Dong, F. Liao, T. Pang, H. Su, J. Zhu, X. Hu, and J. Li, "Boosting adversarial attacks with momentum," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 9185–9193.
- [37] Z. Zhao, H. Zhang, R. Li, R. Sicre, L. Amsaleg, M. Backes, Q. Li, Q. Wang, and C. Shen, "Revisiting transferable adversarial images: Systemization, evaluation, and new insights," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 48, no. 1, pp. 765–780, 2026.
- [38] D. Ravikumar, S. Kodge, I. Garg, and K. Roy, "TREND: Transferability-based robust ENsemble design," *IEEE Transactions on Artificial Intelligence*, vol. 4, no. 3, pp. 534–548, 2023.
- [39] M. Naseer, K. Ranasinghe, S. Khan, F. S. Khan, and F. Porikli, "On improving adversarial transferability of vision transformers," in *International Conference on Learning Representations*, 2022.
- [40] Z. Wei, J. Chen, M. Goldblum, Z. Wu, T. Goldstein, and Y.-G. Jiang, "Towards transferable adversarial attacks on vision transformers," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 36, no. 3, 2022, pp. 2668–2676.
- [41] J. Zhang, Y. Huang, W. Wu, and M. R. Lyu, "Transferable adversarial attacks on vision transformers with token gradient regularization," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2023, pp. 16 415–16 424.
- [42] W. Ma, Y. Li, X. Jia, and W. Xu, "Transferable adversarial attack for both vision transformers and convolutional networks via momentum integrated gradients," in *IEEE/CVF International Conference on Computer Vision*, 2023, pp. 4630–4639.
- [43] Z. Chen, C. Xu, H. Lv, S. Liu, and Y. Ji, "Understanding and improving adversarial transferability of vision transformers and convolutional neural networks," *Information Sciences*, vol. 648, p. 119474, 2023.
- [44] F. Pinto, P. H. S. Torr, and P. K. Dokania, "An impartial take to the CNN vs Transformer robustness contest," in *Computer Vision – ECCV 2022*, ser. Lecture Notes in Computer Science, vol. 13673. Cham: Springer, 2022, pp. 466–480.
- [45] K. Ali, M. S. Bhatti, A. Saeed, A. Athar, M. A. Al Ghamdi, S. H. Almotiri, and S. Akram, "Adversarial robustness of vision transformers versus convolutional neural networks," *IEEE Access*, vol. 12, pp. 105 281–105 293, 2024.
- [46] A. Krizhevsky, G. Hinton *et al.*, "Learning multiple layers of features from tiny images," University of Toronto, Tech. Rep., 2009.
- [47] S. Zagoruyko and N. Komodakis, "Wide residual networks," in *Proceedings of the British Machine Vision Conference (BMVC)*. BMVA Press, 2016, pp. 87.1–87.12.
- [48] S. Goyal, C. Qin, J. Uesato, T. Mann, and P. Kohli, "Uncovering the limits of adversarial training against norm-bounded adversarial examples," *arXiv preprint arXiv:2010.03593*, 2020.
- [49] R. Wightman, "Pytorch image models," <https://github.com/rwightman/pytorch-image-models>, 2019.
- [50] E. Debenedetti, V. Sehwag, and P. Mittal, "A light recipe to train robust vision transformers," in *IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*, 2023, pp. 225–253.
- [51] Y. Mo, D. Wu, Y. Wang, Y. Guo, and Y. Wang, "When adversarial training meets vision transformers: Recipes from training to architecture," in *Advances in Neural Information Processing Systems*, vol. 35, 2022, pp. 18 599–18 611.
- [52] L. Rice, E. Wong, and J. Z. Kolter, "Overfitting in adversarially robust deep learning," in *International Conference on Machine Learning*, 2020, pp. 8093–8104.
- [53] I. Loshchilov and F. Hutter, "SGDR: Stochastic gradient descent with warm restarts," in *International Conference on Learning Representations (ICLR)*, 2017.
- [54] —, "Decoupled weight decay regularization," in *International Conference on Learning Representations (ICLR)*, 2019.
- [55] N. Hurley and S. Rickard, "Comparing measures of sparsity," *IEEE Transactions on Information Theory*, vol. 55, no. 10, pp. 4723–4741, Oct. 2009.
- [56] S.-M. Moosavi-Dezfooli, A. Fawzi, O. Fawzi, and P. Frossard, "Universal adversarial perturbations," in *IEEE Conference on Computer Vision and Pattern Recognition*, 2017, pp. 1765–1773.
- [57] Z. Wang, T. Pang, C. Du, M. Lin, W. Liu, and S. Yan, "Better diffusion models further improve adversarial training," in *Proceedings of the 40th International Conference on Machine Learning (ICML)*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 2023, pp. 36 246–36 263.

PLACE
PHOTO
HERE

Çağrı Şahin received the B.Sc. degree in software engineering from Fırat University, Elazığ, Turkey. He is currently pursuing the M.Sc. degree at the Department of Software Engineering, Fırat University. His research interests include adversarial machine learning, deep learning security, and computer vision, with a focus on the robustness of neural network architectures against adversarial perturbations.

PLACE
PHOTO
HERE

Özal Yıldırım is with the Department of Software Engineering, Fırat University, Elazığ, Turkey. His research interests include deep learning, computer vision, and biomedical signal processing.